



AI Engineer Interview

Deliverable: AI Pipeline

2026

1 Overview

This assignment forms the basis of the evaluation for the next and final stage of our recruitment process. The objective is to deliver an AI solution to a realistic business case, implemented as a pipeline.

The deliverable is a single link to a public GitHub repository.

Should you progress to the next stage, we will discuss your submission and the reasoning behind the decisions you made. You are welcome to use AI assistance for configuration, for code production, and for understanding the technical requirements, provided that you are able to justify every choice in your solution.

We are looking for motivated engineers who enjoy building and architecting systems around AI. There is no single correct answer to this assignment, but we recommend familiarising yourself with the fundamentals of AI engineering, such as evaluation and optimisation, before you begin.

Please include a clear and concise README, so that the reviewer can assess your work without difficulty.

You are free to be creative in how you design your solution, and to draw on your own strengths and the areas you are most interested in.

We look forward to reading your submission. Best of luck.

2 The Case: ORION

ORION (Operational Risk & Integrity Office) is a public authority responsible for approving firms that provide high-risk financial and digital infrastructure services within a regulated economic zone. Firms must obtain authorization before they may operate. Applications vary widely in quality, size, and completeness, and frequently include large volumes of unstructured material.

2.1 Goal

Design and implement a system that reviews authorization submissions and produces a structured risk assessment suitable for review by a human analyst.

2.2 Input Format

Each submission arrives as:

- A primary JSON file describing the applicant and the declared activities.
- References to a document set held in external storage (PDF, DOCX or XLSX).

The document set may exceed 100 pages in total, and may contain redundant or irrelevant material.

2.3 Output Format

The system must produce:

- Risk ratings across predefined dimensions.
- Derived composite scores.
- A recommended authorization level.
- Clarification or follow-up questions where information is missing or inconsistent.
- A structured payload suitable for reviewer validation and override.

2.4 Constraints

- The system executes as a serverless workflow with a fixed execution limit.
- Input metadata and documents are referenced from object storage.
- Results must be delivered to an external review API.
- An LLM may be used for interpretation and reasoning, via an LLM API.
- Outputs must be auditable and reproducible.

2.5 Deliverable

Design a pipeline that:

- Ingests submission metadata and documents.
- Extracts the relevant facts and signals.
- Produces risk ratings and composite scores.
- Determines an authorization level.
- Generates follow-up questions where appropriate.
- Emits a reviewer-ready result.

Provide a minimal but complete implementation, in a GitHub repository that:

- Can be executed locally or in a container.
- Demonstrates ingestion, processing, and output.
- Produces structured, validated results.
- Includes logging or tracing for auditability.